

IronAI: Intelligent Fitness Agent

A Graph-Based LLM Architecture for Hybrid Fitness Coaching

Jhony Peñaherrera

December 4, 2025

Intelligent Agents Course - Final Project

The Problem: Analysis Paralysis in Fitness

Why Traditional Apps Fail

- Too rigid and formulaic
- No personalized motivation
- Poor natural language understanding

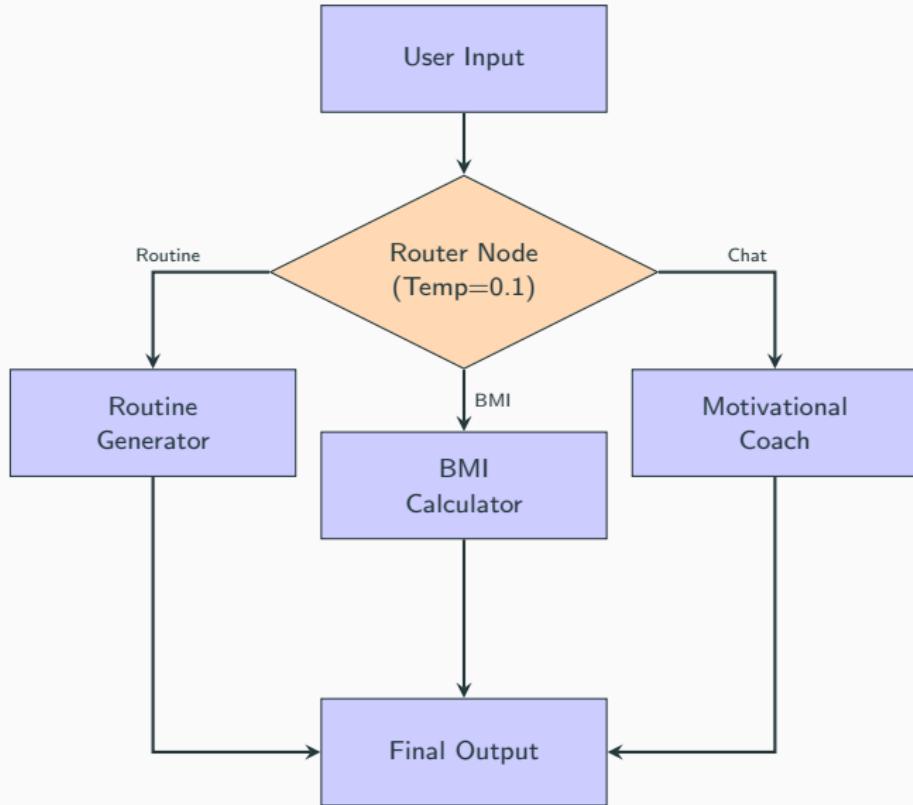
Why Generic LLMs Fail

- Hallucinate health calculations
- Inconsistent advice
- No specialization for fitness

IronAI Solution: Hybrid Intelligence

Combines **strict mathematical precision** (BMI, calories) with **creative motivational coaching** using temperature-controlled routing in a graph architecture.

System Architecture: Graph-Based Routing



Tech Stack: Python, LangChain, LangGraph, Ollama (Gemma 3:1b)

LangGraph: State Definition

Why State Management?

Multi-step workflows require persistent context across nodes

```
1 from typing import TypedDict
2 from langgraph.graph import StateGraph
3
4 class AgentState(TypedDict):
5     input_text: str          # User query
6     intent: str              # Router classification
7     user_data: dict          # Extracted parameters
8     final_output: str        # Agent response
```

- **StateGraph:** Manages transitions between nodes
- **Conditional Edges:** Route based on intent
- **Persistence:** Each node reads/writes to shared state

The Router: Intent Classification with Precision

```
1 def router_node(state: AgentState):
2     prompt = """
3         Classify intent as:
4         - 'routine': workout plans
5         - 'bmi': health calculations
6         - 'chat': motivation/advice
7         """
8
9     llm = ChatOllama(
10        model="gemma3-1b",
11        temperature=0.1    # Precision!
12    )
13
14     intent = llm.invoke(prompt)
15     return {"intent": intent}
```

Temperature Strategy

Router temp=0.1

*Deterministic
classification*

Coach temp=0.7

*Creative
responses*

Key LangGraph Feature:
Conditional edges route to
different nodes based on
`state["intent"]`

Prompt Engineering: Few-Shot & Sequential Chains

Few-Shot Prompting for Personality (Chat Node)

- Define “IronCoach” persona: *Rude but motivating, like a drill sergeant*
- Provide 3-5 examples of typical exchanges in prompt template
- Result: Consistent personality across conversations

Sequential Chain (Routine Node)

Chain 1: Generate 5 exercises for target muscle group

↓ *Output becomes input*

Chain 2: Recommend supplements based on exercises

Implemented using LLMChain with PromptTemplate from LangChain

All prompts isolated in /prompts/templates.py for maintainability

Tool Integration: BMI Calculator Node

Step 1: Extract Parameters

```
1  from langchain.output_parsers
2      import JsonOutputParser
3
4  parser = JsonOutputParser()
5  prompt = """
6      Extract weight and height:
7      Input: "I weigh 80kg and
8          I'm 175cm tall"
9      Output: {"weight": 80,
10             "height": 1.75}
11      """
12
13  data = parser.parse(
14      llm.invoke(prompt)
15  )
```

Step 2: Execute Tool

```
1  # tools.py
2  def calculate_bmi_tool(
3      weight: float,
4      height: float
5  ):
6      bmi = weight / (height ** 2)
7
8      if bmi < 18.5:
9          return "Underweight"
10     elif bmi < 25:
11         return "Normal"
12     else:
13         return "Overweight"
14
15 # Call from node
16 result = calculate_bmi_tool(
17     **data
18 )
```

Critical: `JsonOutputParser` ensures reliable extraction (no regex hacks!)

Production-Ready Code Structure

```
1 ironai/
2
3     main.py
4     requirements.txt
5
6     src/
7         nodes.py
8         graph.py
9         state.py
10        tools.py
11
12    prompts/
13        templates.py
14
15    logs/
16        agent.log
```

Design Principles

- **Separation of Concerns:** Logic, prompts, and tools isolated
- **Testability:** Each node can be unit tested
- **Scalability:** Easy to add new nodes/intents
- **Logging:** Track routing decisions for debugging

Entry point (main.py) compiles graph and runs agent loop

Demonstration: Three Test Cases

1. Routine Generation

- Input: *"Give me a chest routine for hypertrophy"*
- Router → Routine Node (Sequential Chain)
- Output: 5 exercises + supplement recommendations

2. BMI Calculation

- Input: *"Calculate my BMI, I weigh 80kg and I'm 175cm"*
- Router → BMI Node (Parameter Extraction + Tool)
- Output: *"Your BMI is 26.1 (Overweight). Consider caloric deficit."*

3. Motivational Coaching

- Input: *"I'm feeling tired today"*
- Router → Chat Node (Few-Shot Personality)
- Output: *"Tired? Your goals don't care! 10 push-ups NOW!"*

Key Success: Router correctly classifies intent in all cases (temp=0.1)

Conclusion: Why IronAI Works

Technical Achievements

- **LangChain:** Sequential chains, prompt templates, output parsers
- **LangGraph:** StateGraph with conditional routing
- **Prompt Engineering:** Few-shot for personality, structured extraction
- **Tool Integration:** Python function execution via LangChain
- **Temperature Control:** 0.1 (precision) vs 0.7 (creativity)

The Bigger Picture

Local Execution (Ollama) + **Graph Logic** (LangGraph) > Generic Chatbots

IronAI demonstrates that specialized agents with hybrid architectures outperform general-purpose LLMs for domain-specific tasks.